

Natural Language Processing (NLP) with Machine Learning – Master Notes

What is NLP?

Definition

Natural Language Processing (NLP) is a branch of Artificial Intelligence (AI) and Machine Learning (ML) that enables computers to understand, interpret, analyze, and generate human language. It bridges human communication and computer understanding.




Goals of NLP

- ✓ Convert unstructured text structured numeric data
- ✓ Enable ML models for classification, prediction, clustering
- ✓ Capture context, semantics, and syntax

Applications

Application	Example
Text Classification	Spam detection, sentiment analysis
Machine Translation	English French
Summarization	Auto-summarize articles
Chatbots & Assistants	Siri, Alexa, Google Assistant
Question Answering	Search engines, customer support

2 Key Terms

Term	Definition	Example
 Corpus	Collection of text	100 movie reviews
 Document	Single piece of text	Sentence, paragraph, or article
 Vocabulary	All unique words in corpus	"I love pizza" + "I love pasta" {I, love, pizza, pasta}

3 Text Preprocessing

Definition

Preprocessing = Cleaning and standardizing text for ML models.

Steps

- 1 Lowercasing: "I Love NLP" "i love nlp"
- 2 Tokenization: "i love nlp" ["i", "love", "nlp"]
- 3 Stopword Removal: Remove common words ["love", "nlp"]
- 4 Stemming & Lemmatization:
 - ◆ Stemming: "running" "run"
 - ◆ Lemmatization: "better" "good"
- 5 Vectorization: Convert text numeric vectors (see below)

4 Feature Extraction / Vectorization Techniques

1 One-Hot Encoding (OHE)

Definition

Converts text or categorical data into a binary vector. Each unique word gets 1 in its index and 0 elsewhere.

Example

Vocabulary = {I, love, NLP}

"love NLP" [0, 1, 1]

✓ Pros: Simple, easy

✗ Cons: Sparse, no semantics

2 Bag of Words (BoW)

✓ Definition

Represents text as a vector of word counts. Each position = frequency of a word in the document.

Example

Vocabulary = {I, love, NLP, fun}

"I love NLP NLP" [1, 1, 2, 0]

✓ Pros: Simple, fast

✗ Cons: Ignores word order, context, semantics

3 TF-IDF (Term Frequency \u2013 Inverse Document Frequency)

✓ Definition

Weights words by importance: frequent in document but rare in corpus.

$$TF\text{-}IDF(t,d) = TF(t,d) \times \log\frac{N}{DF(t)}$$

TF(t,d): Term frequency in document

DF(t): Number of documents containing term

N: Total documents

✓ Pros: Highlights key words

✗ Cons: Sparse for large vocab

4 Word Embeddings

✓ Definition

Dense vector representations capturing semantic meaning. Words with similar meanings are close in vector space.

Examples: Word2Vec, GloVe, FastText, BERT embeddings

Example

"king" \u2013 "man" + "woman" \u2013 "queen"

✓ Pros: Captures meaning, context

✗ Cons: Requires large corpus

5 N-grams

✓ Definition

Contiguous sequence of n words \u2013 partial context.

Unigram: "I", "love", "NLP"

Bigram: "I love", "love NLP"

Trigram: "I love NLP"

✓ Pros: Adds semantic meaning

✗ Cons: High dimensionality, OOV problem

5 Python Examples

Bag of Words

```
from sklearn.feature_extraction.text import CountVectorizer
docs = ["I love NLP", "NLP is fun"]
vectorizer = CountVectorizer()
X = vectorizer.fit_transform(docs).toarray()
print(vectorizer.get_feature_names_out())
print(X)
```

TF-IDF

```

from sklearn.feature_extraction.text import TfidfVectorizer
docs = ["I love NLP", "NLP is fun"]
vectorizer = TfidfVectorizer()
X = vectorizer.fit_transform(docs).toarray()
print(vectorizer.get_feature_names_out())
print(X)

```

Word2Vec

```

from gensim.models import Word2Vec
sentences = [["I", "love", "NLP"], ["NLP", "is", "fun"]]
model = Word2Vec(sentences, vector_size=5, window=2, min_count=1)
print(model.wv['NLP'])

```

6 Text Classification with ML

✓ Algorithms: Naïve Bayes, Logistic Regression, SVM, Random Forest

✓ Metrics: Accuracy, Precision, Recall, F1-score

\u25c0 Example (TF-IDF + Naïve Bayes):

```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB

docs = ["I love NLP", "NLP is fun", "I hate spam"]
labels = [1, 1, 0]

vectorizer = TfidfVectorizer()
X = vectorizer.fit_transform(docs)

model = MultinomialNB()
model.fit(X, labels)
print(model.predict(vectorizer.transform(["I love spam"])))

```

7 Topic Modeling (Unsupervised)

✓ Algorithms: LDA, NMF

◆ Example:

```
from sklearn.decomposition import LatentDirichletAllocation
from sklearn.feature_extraction.text import CountVectorizer

docs = ["I love NLP", "NLP is fun", "I hate spam"]
vectorizer = CountVectorizer()
X = vectorizer.fit_transform(docs)

lda = LatentDirichletAllocation(n_components=2, random_state=0)
lda.fit(X)
print(lda.components_)
```




8 Sequence Models (RNN, LSTM)






✓ RNN: Maintains hidden state

✓ LSTM: Handles long-term dependencies


✓ Use Cases: Next-word prediction, sentiment analysis, chatbots

9 Quick Comparison Table

Technique	Type	✓ Pros	✗ Cons	Use Cases
 One-Hot Encoding	Feature Extraction	Simple	Sparse, no semantics	Small datasets
 BoW	Feature Extraction	Simple, fast	Ignores context, sparse	Text classification
 TF-IDF	Feature Extraction	Highlights important words	Sparse for large vocab	Document retrieval

 Word Embeddings	Feature Extraction	Captures semantic meaning & context	Requires large corpus	Similarity, chatbots
 N-grams	Feature Extraction	Adds some context	High dimension, OOV	Text classification
 LDA	Unsupervised	Finds hidden topics	Needs preprocessing	Topic modeling
 ML Classifiers	Supervised	Predicts labels	Needs labeled data	Spam detection, sentiment
 RNN/LSTM	Sequence Modeling	Handles sequential dependencies	Computationally expensive	Text generation chatbots



 Now the notes are enhanced with symbols, icons, and clear formatting for better readability and memory retention.